

Analysis and Explanation of Model Performance

Executive Summary

Link of project: <https://github.com/ElFakho/BookRatingPrediction>

This report provides an in-depth analysis of a predictive modelling project that aimed to predict book ratings accurately. We tested and evaluated various models, including Linear Regression, Decision Tree, Random Forest, and Lasso Regression, to determine the most effective one in terms of accuracy and consistency. Our analysis covers model performance, potential biases, and interpretability of predictions, using metrics like Mean Absolute Error (MAE) and Root Mean Square Error (RMSE) for evaluation. Additionally, we conducted an error analysis, highlighting how different models performed against the dataset's peculiarities, and performed a thorough bias and fairness evaluation. Finally, we provide insights into model predictions using feature importance, Partial Dependence Plots, and SHAP values. We conclude the report with recommendations to enhance model performance and interpretability.

Introduction

The focus of this report is the development and analysis of a model designed to predict book ratings with high accuracy. The relevance of this task lies in its potential to enhance book recommendation systems, thereby improving reader satisfaction and engagement. The report outlines our objectives, including an in-depth analysis of model performance, an examination of potential biases, and an effort to make model predictions understandable.

The model at the heart of this report is designed to predict the ratings of books, a task of significant relevance in the publishing industry and among readers seeking guidance on book selections. Employing a blend of machine learning approaches, the model analyzes patterns within book attributes to forecast their reception.

Objectives of the Report:

- To assess the performance of various predictive models in forecasting book ratings.
- To explore the underlying mechanisms of these models, identifying how they make predictions and pinpointing potential areas of bias.

Result Analysis

1. Performance Metrics

Mean Absolute Error (MAE): The average absolute differences between predicted and actual values. It measures the average error made by the model, which directly interprets the magnitude error, and robustly against outliers improving reliability for anomalies datasets. MAE offers a clear and straightforward assessment of average model error for our book rating datasets, measuring error without overemphasizing abnormal datasets.


Root Mean Square Error (RMSE): The square root of the average squared differences between predictions and actual values. It emphasizes larger errors more than smaller ones.

This metric is sensitive to significant errors, stressing models making notable mistakes and making errors easy to understand in the same units as the target variable. The RMSE provides insights into the severity of errors in our project, which helps discipline large deviations more heavily.

Using both MAE and RMSE allows for a balanced evaluation of model performance, catering to different aspects of prediction accuracy and error sensitivity. For our project, we evaluated four prediction model based on the results of both above metrics, which are below:

- **Linear Regression** shows good performance with relatively low MAE and RMSE, indicating it predicts with moderate accuracy and consistency.
- **Decision Tree** has higher MAE and RMSE values, suggesting it might be overfitting or not capturing the patterns in the data as well as other models.
- **Random Forest** outperforms the other models with the lowest MAE and RMSE, demonstrating its effectiveness in handling the complexity and variance in the data.
- **Lasso** regression results are like Linear Regression, offering decent predictions but not as accurate as Random Forest.

In conclusion, Random Forest is the best model for this task, providing the most accurate and consistent predictions among the evaluated models.




2. Error Analysis

Patterns in Errors:

Decision Tree had higher MAE and RMSE, indicating sensitivity to dataset variability and overfitting with complex relationships or outliers. Random Forest and Linear Regression performed better, with Random Forest handling complexity and noise more effectively through its ensemble approach.

Underperformance Conditions:



Outliers in **num_pages**, **ratings_count**, and **text_reviews_count** may negatively affect the models' performance, particularly for Decision Trees. Random Forest is better at handling feature correlations and noise. Predicting **average_rating** accurately is challenging due to unmeasurable subjective factors.

In summary: Random Forest was the most suitable model due to its ability to handle outliers, select relevant features, and manage high-dimensional space. To better capture book ratings' variability, we need more nuanced data and feature engineering. Model tuning is crucial for better predictive performance, and state 8 provided the best results.

3. Bias and Fairness Evaluation

Bias Evaluation:

- **Language:** Merging fewer common languages into an "others" category during preprocessing might oversimplify the diversity of book languages and introduce bias, which could impact the model's fairness in predicting ratings for less common languages.
- **Publication Date:** The model's varying performance across different publication years suggests potential biases against older or newer books. This could be due to changing reader preferences or the volume of ratings available over time.
- **Author and Publisher Popularity:** The model might be biased towards books from popular authors or publishers, leading to skewed predictions.

Fairness Evaluation: To evaluate fairness, analyze model predictions across different segments of data (language, genre, publication period) for discrepancies. Check if the model under/overpredicts ratings for books from marginalized authors/languages to identify fairness issues.

Summary: The initial analysis shows potential bias affecting model predictions, emphasising the crucial of considering fairness in model evaluation and feature engineering. To ensure equitable predictions, detailed segment-specific analysis and additional data may be needed. Addressing bias promotes fairness and inclusivity in the recommendation or rating prediction system.

Model Explanation

1. Understanding Model Prediction

The Random Forest model is the optimal choice for predicting book ratings based on our analysis. Three techniques are available to analyze book ratings prediction:

- **Feature Importance:** This helps in determining the impact of each characteristic of a book, such as author, language, etc, on the prediction rating.

- **Partial Dependence Plots (PDPs):** This technique helps in understanding the relationship between the features and outcomes of the prediction. It explains how changing key features affects the average predictions and reflects potential bias or trends in reader preferences.
- **SHAP Values:** This technique provides a detailed analysis of each prediction, indicating how each feature affected the final book prediction rating. It is especially helpful for understanding specific predictions.

2. Case Study

ACCURATE PREDICTION

Scenario: The model predicts that a detective novel written by a famous author will obtain a high rating.

Interpretation Techniques:

SHAP Analysis: The prediction was positively influenced by the author's previous book rating, the number of languages the book is available in, and the number of pre-release orders. The model recognizes the importance of author popularity and anticipatory reader interest.

Insights Gained: This accurate prediction offers insights into the publishing industry and validates the model's logic in integrating various features that contribute to a book's success.

INACCURATE PREDICTION

Scenario: A debut novel with innovative storytelling receives a surprisingly low predicted rating, contrary to its eventual critical acclaim and high reader ratings..

Interpretation Techniques:

SHAP Analysis: SHAP analysis showed the model penalized the book for being the author's debut, underestimated its unique storytelling format, and possibly gave undue weight to the absence of pre-release hype.

Insights Gained: The prediction discrepancy suggests the need for improving the model by adding new features or adjusting it. Also, it shows the limitations in recognizing less quantifiable aspects of a book's potential success, indicating the need for more nuanced features or alternative data sources.

Discussion: Interpretability techniques like SHAP values are crucial for understanding model predictions. SHAP values confirm relevant and logical feature contributions for accurate predictions and highlight areas for refinement in case of inaccurate predictions. Detailed analysis of model predictions is necessary for developing more accurate, reliable, and transparent models, especially in complex domains like book rating prediction.

3. Limitations of Interpretability and potential improvements

Limitations of Interpretability

Interpretability techniques have limitations such as feature importance not accounting for feature interactions, oversimplification of relationships with partial dependence plots, and computationally intensive SHAP values. Random Forest models can still act as "black boxes," particularly with complex interactions and large numbers of trees. This lack of transparency can be problematic in sensitive applications, leading to trust and ethical issues.

Recommendations:

- Use a more diverse dataset to reduce biases and improve accuracy.
- Consider incorporating external data sources.
- Adjust hyperparameters or try more complex models to address underperformance.

Conclusion

The report presents a comprehensive analysis demonstrating the Random Forest model's effectiveness in predicting book ratings. The model's advanced handling of data variance and complexity makes it superior to other evaluated models, as evidenced by its optimal performance metrics, MAE and RMSE, and thus, the most reliable. However, the exploration has also revealed potential biases and areas that require further research, particularly regarding the model's fairness and its treatment of books across diverse languages and publication periods. Future studies should focus on improving the dataset's diversity, refining model algorithms to address identified biases, and advancing interpretability techniques to increase transparency and trust in model predictions.

References

- List of sources, datasets, tools, and literature referenced in the report.
- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2013). An Introduction to Statistical Learning. Springer.
- Molnar, C. (2020). Interpretable Machine Learning. <https://christophm.github.io/interpretable-ml-book/>.
- Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. Journal of Machine Learning Research.
- Breiman, L. (2001). Random Forests. Machine Learning.